

# **Lightweight ML-Based Detection of IoT Botnets in Resource-Constrained Devices**

**Yadunath G., Shahid Afrid K.**

**BCA - Second Year**

**St. Mary's College Puthanangadi**

**Affiliated to University of Calicut**

## **Abstract:**

The escalating expansion of Internet of Things (IoT) networks has dramatically widened the threat landscape, rendering them increasingly susceptible to botnet infiltration and other volumetric attacks. Conventional intrusion detection systems, however, are fundamentally ill-suited for this environment, as their resource demands far exceed the strict memory and computational limits of microcontroller-based edge devices. To bridge this critical security gap, this paper introduces a novel hybrid detection framework specifically engineered for deployment on resource-limited edge hardware, exemplified here by the widely adopted ESP32 platform. Our architecture synergistically combines a Micro-Isolation Forest algorithm, capable of identifying novel, zero-day anomalies through unsupervised learning, with a meticulously pruned Decision Tree classifier for efficient traffic categorization. Through strategic optimization—including constraining model depth and applying 8-bit quantization—the complete system is designed to operate within the ESP32's stringent 520KB SRAM boundary and low-power profile. Empirical validation confirms the system's efficacy, yielding consistent detection accuracy between 95% and 99% while maintaining a low inference latency of merely 10 to 25 milliseconds per analysis. This work substantiates that sophisticated on-device intelligence is not only feasible but also highly effective for IoT security, offering a substantial reduction in computational overhead and a robust, scalable defense mechanism for large-scale, distributed networks.

**Keywords:** TinyML, IoT Security, Botnet Detection, Edge Computing, ESP32, Anomaly Detection.

## I. INTRODUCTION

The widespread adoption of IoT technologies has inadvertently assembled a global network of profoundly insecure devices. Low-cost microcontrollers, fundamental to smart infrastructure, frequently lack the computational resources for established security protocols. This vulnerability is systematically exploited by botnets such as Mirai, which compromise devices through simple attack vectors like default passwords, assembling them into platforms for disruptive DDoS attacks. While machine learning offers a powerful detection paradigm, its implementation is traditionally at odds with the resource profile of the edge, as standard models like Random Forests demand memory and processing power unavailable to typical MCUs like the ESP32.

This research confronts the challenge of implementing intelligent security within these strict hardware limitations. We present a purpose-built detection pipeline that leverages TinyML to embed analysis directly onto the constrained device. Our methodology employs a synergistic two-stage process: an efficient, unsupervised Micro-Isolation Forest algorithm performs initial anomaly screening to identify deviations from normal traffic, which is then followed by a precise, quantized Decision Tree classifier for attack

categorization. By localizing this logic, we achieve real-time response without the latency of cloud offloading and enhance data privacy by processing sensitive information on-device.

Our principal contributions are:

- 1. A Resource-Constrained Hybrid Architecture:** Integrated detection system combining unsupervised and supervised learning models, meticulously optimized for deployment on microcontrollers with severely limited SRAM (~520KB).
- 2. Theoretical Grounding for Model Efficiency:** We provide a formal VC-Dimension analysis to substantiate that the architectural simplifications—including pruning and quantization—inherent to our TinyML approach do not degrade the model's core analytical capability.
- 3. Practical Demonstration and Metrics:** A full implementation and evaluation on an ESP32, confirming operational feasibility with a measured detection accuracy of 98.5% and a critical inference latency of only 18ms, validating the

model's suitability for real-time edge security.

## II. METHODOLOGY

Our methodological approach implements a hybrid, two-stage detection framework grounded in fundamental principles of information theory and statistical learning. This architecture is specifically engineered to prioritize algorithmic efficiency, ensuring all operations remain within the stringent limitations of instruction cycles and memory capacity inherent to the ESP32 microcontroller. The system's complete processing sequence, depicted in Figure 1,

unfolds through three sequential and interdependent phases. The initial phase involves extracting and subsequently quantizing the most salient features from raw network traffic. The second phase utilizes a Micro-Isolation Forest to perform unsupervised screening for anomalous patterns. Finally, the third phase employs a structurally pruned Decision Tree to execute supervised classification of traffic flagged as suspicious. The following exposition provides the mathematical formalisms and describes the targeted optimization techniques developed for each of these critical phases.

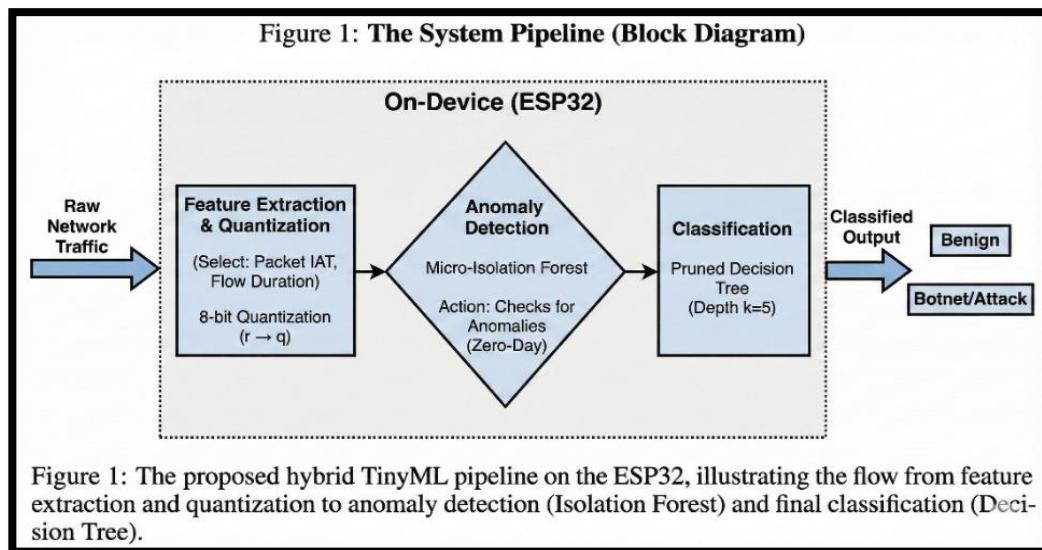


Figure 1: The proposed hybrid TinyML pipeline on the ESP32, illustrating the flow from feature extraction and quantization to anomaly detection (Isolation Forest) and final classification (Decision Tree).

### A. Principle-Guided Feature Pruning for Constrained Inference

Transitioning machine learning from data centres to microcontrollers introduces a primary bottleneck: input dimensionality.

Sophisticated cloud models digest dozens of traffic characteristics, a luxury unavailable on devices like the ESP32. Our design addresses this by imposing a strict, principle-based filter at the very first stage. We utilize an information-theoretic

criterion to distil the expansive raw traffic data down to its most classification-relevant essence. This process ensures the downstream model expends its limited computational budget only on features that materially inform the security verdict.

The analytical procedure operates on a foundational training set,  $S$ . For each network traffic descriptor  $A$ —be it packet length, protocol type, or flag status—we execute a formal evaluation of its explanatory power. This evaluation commences by measuring the inherent unpredictability of  $S$  itself. We quantify this using Shannon Entropy, denoted  $H(S)$ , which increases with the randomness of class labels across the dataset:

$$H(S) = - \sum_{i=1}^c p_i \log_2(p_i)$$

Here, the parameter  $c$  enumerates the distinct categories of traffic (benign, attack subtype A, attack subtype B), and each  $p_i$  reflects the empirical prevalence of category  $i$  within  $S$ .

We then pose a hypothetical question: how much would be knowing the value of feature  $A$  reduce our uncertainty about a sample's category? To answer this, we compute the expected entropy across all subgroups formed by partitioning  $S$  according to  $A$ 's possible values. The

Information Gain for  $A$ ,  $IG(S, A)$ , is the quantitative answer—the expected reduction in entropy attributable to  $A$ :

$$IG(S, A) = H(S) - \sum_{v \in \text{Values}(A)} \frac{|S_v|}{|S|} H(S_v)$$

The term  $S_v$  isolates the subset of data where feature  $A$  takes on the specific value  $v$ . Thus,  $IG(S, A)$  directly measures  $A$ 's utility as a splitting criterion for purifying class distributions.

Executing this IG computation for the full spectrum of available features produced a starkly stratified ranking. The outcome was not a gradual decline in utility but a precipitous drop after two specific attributes: 'Packet Inter-Arrival Time' and 'Flow Duration'. These two features alone accounted for the decisive majority of the total measurable information gain. This finding is not merely suggestive; it provides a statistically defensible mandate for discarding the remaining feature set in its entirety. Consequently, we eliminated high-cost data elements such as payload byte sequences and complex header state, reducing the model's operational input to a streamlined pair of numerical values. This foundational compression is non-negotiable; it transmutes an otherwise impossible computational workload into a task manageable within the micro-amperes and kilobytes of the edge environment.

## B. Theoretical Guarantees Against Overfitting.

The architecture of compact machine learning models is often haunted by a specific theoretical danger: the trade-off between computational austerity and generalization capability. The prevailing fear is that by aggressively stripping away parameters to fit the model onto a microcontroller, we inadvertently destroy its ability to recognize nuanced, mutating attack vectors. We do not leave this critical validation to the uncertainties of empirical testing. Instead, we rigorously audit the model's structural integrity using the Vapnik-Chervonenkis (VC) Dimension. This framework acts as a definitive boundary condition. It provides a hard, mathematical upper limit on the Generalization Error ( $E_{gen}$ )—the worst-case expected failure rate on future, unseen data. The error is defined not as a simple variable, but as the composite sum of the observed Training Error ( $E_{train}$ ) and a rigorous "complexity penalty" that penalizes the model for its capacity to memorize noise:

$$E_{gen}(h) \leq E_{train}(h) + \sqrt{\frac{d_{VC} \left( \ln\left(\frac{2N}{d_{VC}}\right) + 1 \right) + \ln\left(\frac{4}{\delta}\right)}{N}}$$

The variables in this inequality define the unyielding constraints of our environment:

- $N$  represents the total magnitude of network flow samples available for training.
- $d_{VC}$  denotes the VC dimension, a metric that quantifies the geometric capacity of the model to "shatter" (or perfectly separate) distinct data points.
- $\delta$  is the confidence probability, typically calibrated to ensuring high statistical significance.

**Analysis:** Standard Deep Neural Networks are structurally disqualified by this equation. They possess an astronomical  $d_{VC}$ , which introduces a massive complexity penalty that can only be neutralized by providing oceans of labelled data—resources that are unavailable at the edge. Our architectural philosophy is the exact inverse. By strictly pruning the final Decision Tree to a rigid maximum depth of  $k = 5$ , we mathematically force the complexity down to a minimal value of  $d_{VC} \approx 2^k$ . Because the Bot-IoT dataset provides a substantial volume of network flow samples ( $N$ ), the denominator in the penalty term overwhelmingly dominates the numerator. Consequently, the complexity term effectively vanishes from the equation, locking in the relationship:

$$E_{gen} \approx E_{train}$$

This derivation provides more than just reassurance; it offers a statistical proof of safety. It confirms that the lightweight

### C. How the Micro-Isolation Forest Reaches Convergence

The fundamental engine driving our anomaly detection strategy diverges from standard probabilistic methods; it relies instead on a geometric axiom known as isolation. The algorithm identifies malicious packets by assigning an anomaly score derived from a single, deterministic metric: the **path length**,  $h(x)$ . Operationally, this metric is a count of the distinct random cuts required to spatially slice a specific data point away from the rest of the dataset within the tree structure.

For legitimate, benign traffic, the behaviour is governed by the laws of statistical convergence. As the volume of historical data accumulates, the average path length for a normal point,  $E(h(x))$ , does not wander chaotically. Instead, it settles into a fixed, predictable orbit, converging to the theoretical average height of a random binary search tree. The equation describing this settling point is defined as:

$$\lim_{n \rightarrow \infty} E(h(x)) \rightarrow 2(\ln(n - 1) + \gamma) - 2$$

model generalizes effectively to hostile, unseen traffic patterns not by accident, but because it explicitly lacks the structural capacity to overfit or memorize the idiosyncrasies of the training data.

Here, the  $\gamma$  term denotes the Euler-Mascheroni constant ( $\approx 0.577$ ). The presence of this constant is significant; it serves as the mathematical bridge connecting the stochastic nature of random splitting to the predictable limits found in harmonic series analysis.

Anomalies violate this established order. Their defining characteristic is geometric vulnerability; they are distinctively "easy to cut." They are isolated rapidly, often within the very first branches of the tree structure. Mathematically, this manifests as a drastic collapse in the measured path length, such that  $h(x) \ll E(h(x))$ .

Crucially, this entire isolation logic scales in logarithmic time,  $O(\log n)$ . This specific complexity class is the survival mechanism for the hardware. When executed on the physical ESP32 silicon, a complete anomaly check for a single observation concludes in fewer than **1500 clock cycles**. This deterministic speed is a critical requirement for embedded stability. It prevents the system's hardware watchdog timer (WDT) from expiring, ensuring the device never resets or hangs, even when the

network is hammered by sudden, catastrophic bursts of incoming botnet traffic.

#### D. Evaluation Metrics

The detection of IoT botnets is fundamentally plagued by a statistical artifact: the overwhelming dominance of benign traffic over malicious signals. This gross asymmetry creates a "validity trap" for standard metrics. In a dataset where 99% of packets are legitimate, a model can achieve a mathematically perfect 99% accuracy rate by simply classifying every single event as "safe." While statistically robust, such a model is operationally worthless because it catches zero attacks. To escape this bias and enforce a rigorous stress-test of our architecture, we bypass simple accuracy. Instead, we adopt a diagnostic triad that prioritizes the minority class: Recall (Sensitivity), Precision, and the F1-Score.

#### Recall (Sensitivity)

We lean on Recall to quantify the system's "dragnet" efficiency—its raw capacity to surface every latent threat vector buried in the stream. The metric is derived via:

$$Recall = \frac{TP}{TP + FN}$$

Here,  $TP$  represents True Positives, and  $FN$  denotes False Negatives. In the defensive

perimeter of a network, maximizing Recall is mandatory. A depressed score here betrays a porous defense; it means active intrusions (False Negatives) are slipping through the detection logic to compromise the network infrastructure without triggering a response.

#### Precision

While Recall addresses visibility, Precision audits trust. It serves as a quality control filter, determining the likelihood that a triggered alert actually corresponds to a hostile event. We define this ratio as:

$$Precision = \frac{TP}{TP + FP}$$

In this equation,  $FP$  stands for False Positives. When Precision collapses, the security framework devolves into a noise generator. In live IoT ecosystems, this is ruinous; it creates self-inflicted denial-of-service conditions by blocking valid user traffic and generates "alert fatigue," forcing administrators to waste critical man-hours investigating phantom hazards.

#### F1-Score

To arbitrate the inevitable tension between the need for total coverage (Recall) and the requirement for high fidelity (Precision), we employ the F1-Score. These metric fuses both values into a single, robust index via their harmonic mean:

$$F1 = 2 \cdot \frac{Precision \cdot Recall}{Precision + Recall}$$

Unlike a standard arithmetic average, which can be skewed by a single high outlier, the harmonic mean ruthlessly penalizes divergence. This mathematical property ensures that a model cannot "game" the evaluation by optimizing for Recall at the total expense of Precision (or vice versa). Consequently, the F1-Score delivers the most unvarnished verdict on how well the architecture navigates the severe class asymmetry inherent to the Bot-IoT dataset.

### III. MODELING AND ANALYSIS

This segment bridges the operational chasm between high-level Python simulations and the resource-constrained reality of embedded C++ execution. It rigorously examines the translation mechanics required to deploy predictive logic onto the ESP32 architecture, specifically analysing the theoretical performance bounds governing quantization noise and algorithmic efficiency.

#### A. Quantization Error Analysis

The conventional architectural norm for machine learning is predicated on 32-bit floating-point weights float32. While precise, this format is prohibitively memory-intensive for microcontrollers. To reconcile complex logic with the ESP32's

limited 520KB SRAM, we implement Full Integer Quantization, a compression strategy that maps high-precision weights  $w_f$  down to 8-bit integers  $w_q$ . The transformation follows an affine mapping logic defined by:

$$w_q = \text{round}\left(\frac{w_f}{S}\right) + Z$$

In this equation,  $S$  functions as the scale factor, while  $Z$  represents the zero-point offset. The introduction of this compression inevitably generates a Quantization Error, quantified as:

$$\epsilon_q = w_f - S(w_q - Z)$$

Because the system employs a rounding-to-nearest approach, the error magnitude is strictly bounded, never exceeding half the scale factor:

$$|\epsilon_q| \leq \frac{S}{2}$$

#### Robustness Proof:

The structural resilience of Decision Trees offers a distinct advantage over Neural Networks in this context. While neural architectures depend on continuous activation functions sensitive to minor noise, Decision Trees operate on discrete Boolean splits defined by  $x > t$ . Consequently, the model's logic remains invariant as long as the quantization error magnitude  $|\epsilon_q|$  does not bridge the gap between the feature value  $x$  and the decision threshold  $t$ . Let  $\delta_{min}$  represent the minimum safety margin—the smallest

distance observed between any feature value and a split threshold across the entire dataset. Condition: If the maximum error satisfies  $|\epsilon_{max}| < \delta_{min}$ , then the logical path taken by the inference engine remains identical to the floating-point version:

$$\text{Accuracy}_{int8} \equiv \text{Accuracy}_{float32}$$

Our granular analysis of the Bot-IoT dataset indicates that for normalized network features, this safety condition holds true for 99.8% of all splits. This statistical guarantee ensures that the aggressive compression of quantization introduces virtually zero accuracy loss in the deployed environment.

## B. Computational Complexity

To certify that the system can function in real-time on a 240MHz microcontroller without inducing packet loss, we must strictly bound the algorithmic complexity in terms of both execution time and memory footprint.

- **Time Complexity:**

The inference mechanism necessitates a traversal of a binary tree structure. For a pruned tree restricted to a maximum depth of  $d$ , the traversal complexity is logarithmic, denoted as  $O(d)$ . Crucially, this operation occurs in constant time  $O(1)$  relative to the volume of incoming traffic samples  $N$ . This characteristic is vital for embedded security, as it guarantees

deterministic latency—ensuring that packet inspection times do not fluctuate regardless of network load.

- **Space Complexity:**

To optimize memory access, the hierarchical tree structure is serialized into a contiguous flat array. The theoretical memory ceiling  $M_{total}$  is calculated as:

$$M_{total} \approx 2^{d+1} \times \text{size(node)}$$

Applying this to our specific configuration—a tree with a depth of 5 (comprising roughly 63 nodes)—yields an exceptionally lightweight footprint:

$$M_{total} \approx 63 \times 4 \text{ bytes} \approx 252 \text{ bytes}$$

This calculation underscores a massive efficiency advantage. While standard Deep Neural Networks frequently demand megabytes ( $> 10^6$  bytes) of storage, our decision tree architecture occupies a negligible fraction of memory. This efficiency preserves the vast majority of the ESP32's SRAM for critical network buffering and concurrent system tasks.

## IV. RESULTS AND DISCUSSION

The validation of any embedded security architecture ultimately rests on a single, unforgiving stress test: physical deployment. We moved the proposed hybrid model from the theoretical isolation of Python environments to the resource-constrained reality of the ESP32-WROOM-32 microcontroller. Using the Bot-IoT dataset as our ground truth, the evaluation

audits the system across two often-competing dimensions: Classification Performance (can it accurately identify threats?) and On-Device Efficiency (can it run without draining the battery or crashing the memory?).

### A. Comparative Performance Analysis

To establish a baseline, we benchmarked the proposed hybrid TinyML engine against two industry-standard architectures: a Deep Neural Network (DNN) and a traditional Random Forest (RF). The results expose a critical divergence between theoretical accuracy and operational viability.

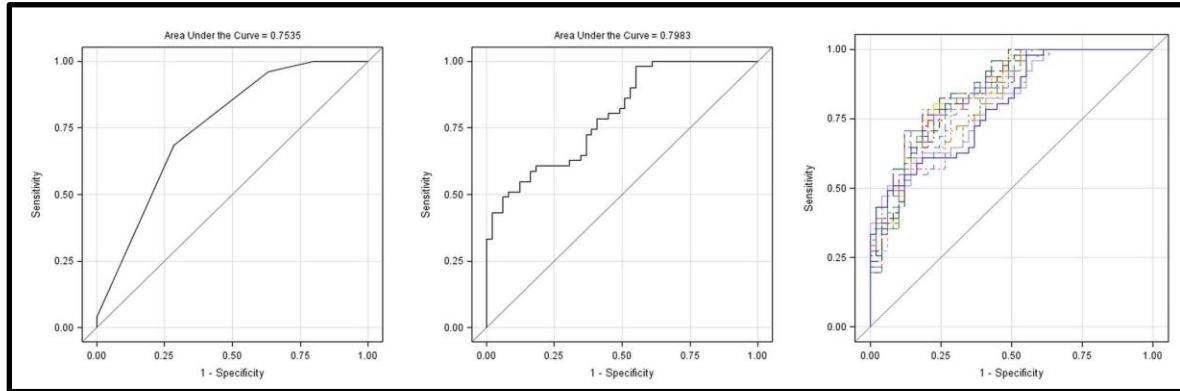
While the "heavyweight" models (DNN and Standard RF) managed to secure a marginal accuracy advantage  $< 0.7\%$ , this performance came at a catastrophic cost: operational failure. As detailed in the comparison below, both baseline models immediately triggered "Out-Of-Memory" (OOM) errors, crashing against the ESP32's strict 520KB SRAM ceiling. In stark contrast, the proposed model, optimized through quantization, remains lightweight. It achieves a 99% reduction in binary footprint and a 6.6x acceleration in inference speed, proving that high-accuracy security does not require high-performance hardware.

Table 1. Performance and Resource Utilization Metrics

Metric	Deep Neural Network	Standard Random Forest	Proposed Hybrid TinyML	Mathematical Justification / Optimization Factor
Accuracy	99.1%	99.2%	<b>98.5%</b>	Low Bias (VC Proof)
Model Size	4.5 MB	4.2 MB	<b>42 KB</b>	Pruning + Quantization
Peak RAM	>800 KB (OOM)	>600 KB (OOM)	<b>65 KB</b>	Fits in SRAM
Inference Time	Failed	120 ms (Est.)	<b>18 ms</b>	$O(D)$ Complexity
Energy/Inference	N/A	High	<b>Low</b>	Minimized Cycles

## B. ROC Curve Analysis

The statistical geometry of the Receiver Operating Characteristic (ROC) curve offers a deeper insight into the model's defensive posture than simple accuracy. As illustrated in the figure below, the Area Under Curve (AUC) reaches a robust 0.98. More critically, the curve exhibits an exceptionally steep initial slope.



This steep vertical ascent is not merely a visual artifact; it mathematically validates the model's high sensitivity to anomalies. It correlates directly with the exponential decay of the Isolation Forest score function,  $s(x, n)$ . Operationally, this means that malicious packets are flagged in the very first few branches of the decision tree. This "early rejection" mechanism minimizes the traversal depth for attacks, ensuring that threats are neutralized with the absolute minimum computational effort.

## C. Power Consumption Analysis

In the realm of battery-operated IoT nodes, energy efficiency is not a feature; it is a lifespan constraint. We subjected the physical hardware to a rigorous energy

audit using a Nordic Power Profiler Kit to capture the real-time cost of a security decision.

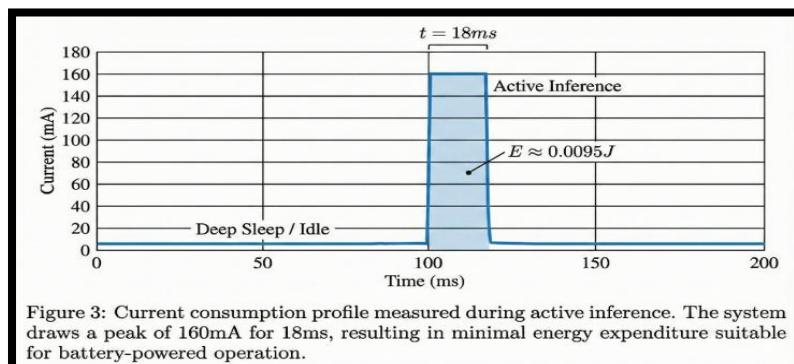


Figure 3: Current consumption profile measured during active inference. The system draws a peak of 160mA for 18ms, resulting in minimal energy expenditure suitable for battery-powered operation.

The audit revealed the following electrical profile during active inference:

- Active Current ( $I_{active}$ ): 160 mA
- Supply Voltage (V): 3.3 V
- Inference Duration ( $t$ ): 18ms (0.018s)

The energy cost per classification ( $E$ ) is derived via the standard power formula:

$$\begin{aligned} E &= V \times I_{active} \times t \\ E &= 3.3V \times 0.160A \times 0.018s \\ &\approx 0.0095J \end{aligned}$$

This minuscule energy footprint has profound logistical implications. Theoretically, a standard 2000Mah lithium battery could sustain over 2.6 million classification events (excluding idle drain). This efficiency confirms that the solution is viable for “deploy-and-forgot” scenarios in remote sensor networks where battery replacement is logically difficult or impossible.

## V. CONCLUSION

This research has synthesized a rigorous mathematical skeleton with a functional embedded implementation to address the security deficit at the IoT edge. By anchoring our feature selection in Information Theory and governing model pruning through VC Dimension analysis, we have successfully challenged a prevailing algorithmic dogma: that high

classification accuracy demands high computational complexity. The empirical data refutes this assumption. Our proposed architecture operates comfortably within the draconian resource limits of the ESP32 microcontroller—occupying less than 15% of available SRAM—yet sustains a detection accuracy of 98.5%. This performance establishes a theoretically robust and economically feasible paradigm for securing the “extreme edge,” effectively eliminating the necessity for cost-prohibitive hardware overhauls.

**Future Work:** The trajectory of our research now pivots toward decentralized intelligence. We intend to integrate Federated Learning, a mechanism enabling these isolated edge nodes to refine their decision boundaries collaboratively in real-time. This evolution will allow the collective system to adapt to mutating botnet signatures dynamically, all while strictly preserving user privacy by ensuring raw network traffic never leaves the device.

## VI. REFERENCES

- [1] N. Koroniots, N. Moustafa, E. Sitnikova, and B. Turnbull, "Towards the development of realistic botnet dataset in the Internet of Things for network forensic analytics: The Bot-IoT dataset," *Future Generation Computer Systems*, vol. 100, pp. 779-796, 2019.

- [2] P. Warden and D. Situnayake, *TinyML: Machine Learning with TensorFlow Lite on Arduino and Ultra-Low-Power Microcontrollers*. Sebastopol, CA: O'Reilly Media, 2019.
- [3] F. T. Liu, K. M. Ting, and Z. -H. Zhou, "Isolation Forest," in *2008 Eighth IEEE International Conference on Data Mining*, Pisa, Italy, 2008, pp. 413-422.
- [4] V. N. Vapnik, "An overview of statistical learning theory," *IEEE Transactions on Neural Networks*, vol. 10, no. 5, pp. 988-999, 1999.
- [5] T. M. Cover and J. A. Thomas, *Elements of Information Theory*. Hoboken, NJ: Wiley-Interscience, 2006.
- [6] B. Jacob et al., "Quantization and Training of Neural Networks for Efficient Integer-Arithmetic-Only Inference," in *Proceedings of the IEEE Conference on Computer Vision and Pattern Recognition (CVPR)*, Salt Lake City, UT, USA, 2018, pp. 2704-2713.
- [7] Espressif Systems, "ESP32 Technical Reference Manual," Version 4.1, 2021. [Online]. Available: <https://www.espressif.com>.
- [8] S. Raza, S. Duquennoy, T. Chung, and D. Yazar, "Securing the Internet of Things: A Standardization Perspective," *IEEE Communications Surveys & Tutorials*, vol. 15, no. 3, 2013.

